

# Amazon Stock Prediction with Linear Regression, LSTM and Prophet

```
In [ ]: import numpy as np
import pandas as pd
import warnings
warnings.filterwarnings('ignore')

import os
for dirname, _, filenames in os.walk('/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

## Read Data

```
In [2]: df=pd.read_csv('../input/nyse/prices.csv')
```

## Describe the data

```
In [3]: df.head()
```

```
Out[3]:      date   symbol     open     close      low      high    volume
0 2016-01-05 00:00:00    WLTW  123.430000  125.839996  122.309998  126.250000  2163600.0
1 2016-01-06 00:00:00    WLTW  125.239998  119.980003  119.940002  125.540001  2386400.0
2 2016-01-07 00:00:00    WLTW  116.379997  114.949997  114.930000  119.739998  2489500.0
3 2016-01-08 00:00:00    WLTW  115.480003  116.620003  113.500000  117.440002  2006300.0
4 2016-01-11 00:00:00    WLTW  117.010002  114.970001  114.089996  117.330002  1408600.0
```

# Check for Null Values

```
In [4]: df.isnull().sum()
```

```
Out[4]: date      0
symbol    0
open      0
close     0
low       0
high      0
volume    0
dtype: int64
```

**No Null values in the entire dataset**

# Check Datatype of the features

```
In [5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 851264 entries, 0 to 851263
Data columns (total 7 columns):
 #   Column   Non-Null Count   Dtype  
 ---  -- 
 0   date      851264 non-null   object 
 1   symbol    851264 non-null   object 
 2   open      851264 non-null   float64
 3   close     851264 non-null   float64
 4   low       851264 non-null   float64
 5   high      851264 non-null   float64
 6   volume    851264 non-null   float64
dtypes: float64(5), object(2)
memory usage: 45.5+ MB
```

```
In [6]: df.describe()
```

Out[6]:

	open	close	low	high	volume
<b>count</b>	851264.000000	851264.000000	851264.000000	851264.000000	8.512640e+05
<b>mean</b>	70.836986	70.857109	70.118414	71.543476	5.415113e+06
<b>std</b>	83.695876	83.689686	82.877294	84.465504	1.249468e+07
<b>min</b>	0.850000	0.860000	0.830000	0.880000	0.000000e+00
<b>25%</b>	33.840000	33.849998	33.480000	34.189999	1.221500e+06
<b>50%</b>	52.770000	52.799999	52.230000	53.310001	2.476250e+06
<b>75%</b>	79.879997	79.889999	79.110001	80.610001	5.222500e+06
<b>max</b>	1584.439941	1578.130005	1549.939941	1600.930054	8.596434e+08

## Dataframe shape

In [7]: `df.shape`

Out[7]: `(851264, 7)`

In [8]: `df['symbol'].nunique()`

Out[8]: `501`

## Extract AMAZON data from the dataset

In [9]: `df1=df[df['symbol']=='AMZN']`

In [10]: `df1.head()`

```
Out[10]:
```

	date	symbol	open	close	low	high	volume
284	2010-01-04	AMZN	136.250000	133.899994	133.139999	136.610001	7599900.0
751	2010-01-05	AMZN	133.429993	134.690002	131.809998	135.479996	8851900.0
1219	2010-01-06	AMZN	134.600006	132.250000	131.649994	134.729996	7178800.0
1687	2010-01-07	AMZN	132.009995	130.000000	128.800003	132.320007	11030200.0
2155	2010-01-08	AMZN	130.559998	133.520004	129.029999	133.679993	9830500.0

```
In [11]:
```

```
df1.describe()
```

```
Out[11]:
```

	open	close	low	high	volume
count	1762.000000	1762.000000	1762.000000	1762.000000	1.762000e+03
mean	337.875664	337.899058	333.969688	341.464438	4.607596e+06
std	189.294231	189.109339	187.654696	190.525796	3.091557e+06
min	105.930000	108.610001	105.800003	111.290001	9.844000e+05
25%	192.962494	193.377506	190.284997	195.532501	2.741550e+06
50%	282.500000	282.915008	279.869995	285.074997	3.890700e+06
75%	398.425003	398.014999	393.799988	402.082496	5.384450e+06
max	845.789978	844.359985	840.599976	847.210022	4.242110e+07

```
In [12]:
```

```
np.round(df1.median(),2)
```

```
Out[12]:
```

```
open      282.50
close     282.92
low       279.87
high      285.07
volume    3890700.00
dtype: float64
```

\*\*\*All the Median of all the numeric values are lesser than mean, so the dataset is right skewed

```
In [13]:
```

```
df1['date']=pd.to_datetime(df1['date'])
```

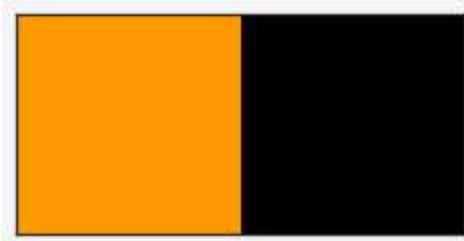
```
In [14]: print("Minimum date value : {}".format(df1['date'].min()))
print("Maximum date value : {}".format(df1['date'].max()))
```

```
Minimum date value : 2010-01-04 00:00:00
Maximum date value : 2016-12-30 00:00:00
```

***we using around 6 years of data***

## EDA

```
In [15]: #importing plotting libraries
import matplotlib.pyplot as plt
import matplotlib.dates as mdates
import seaborn as sns
colors = ['#FF9900', '#000000']
sns.set(palette=colors, font='Serif', style='white', rc={'axes.facecolor':'whitesmoke', 'figure.facecolor':'whitesmoke'}
sns.palplot(colors, size=2)
```

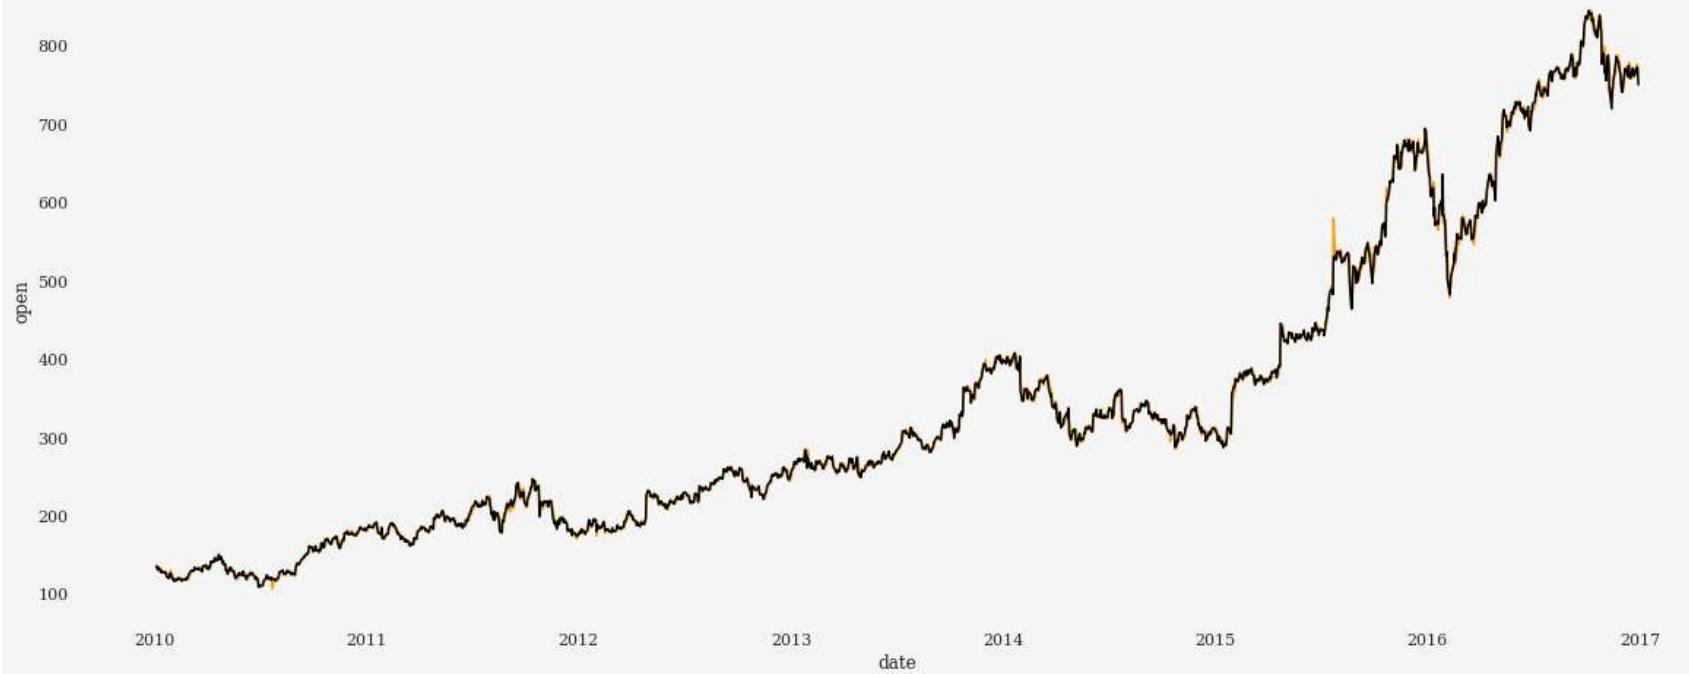


```
In [16]: fig=plt.figure(figsize=(20,8))
ax=sns.lineplot(data=df1, x='date',y='open')
ax=sns.lineplot(data=df1, x='date',y='close', color=colors[1]);
for s in ['left','right','top','bottom']:
    ax.spines[s].set_visible(False)

plt.title("AMAZON Stock value changes since 2010", size=20, weight='bold')
```

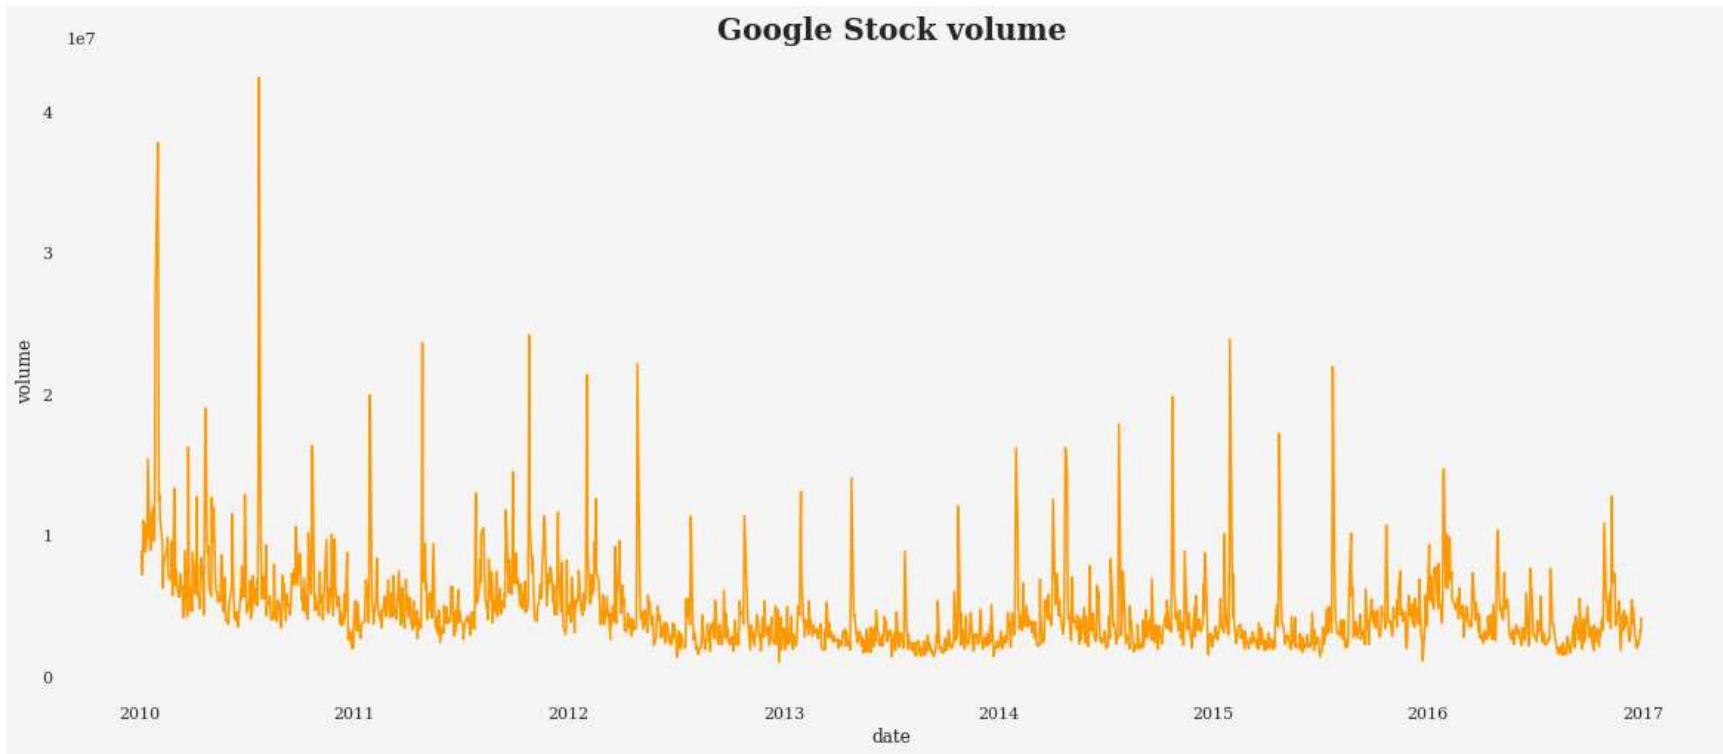
```
Out[16]: Text(0.5, 1.0, 'AMAZON Stock value changes since 2010')
```

## AMAZON Stock value changes since 2010



```
In [17]: fig=plt.figure(figsize=(20,8))
ax=sns.lineplot(data=df1, x='date',y='volume')
#ax=sns.lineplot(data=df1, x='date',y='close', color=colors[1]);
for s in ['left','right','top','bottom']:
    ax.spines[s].set_visible(False)
plt.title("Google Stock volume", size=20, weight='bold')
```

```
Out[17]: Text(0.5, 1.0, 'Google Stock volume')
```



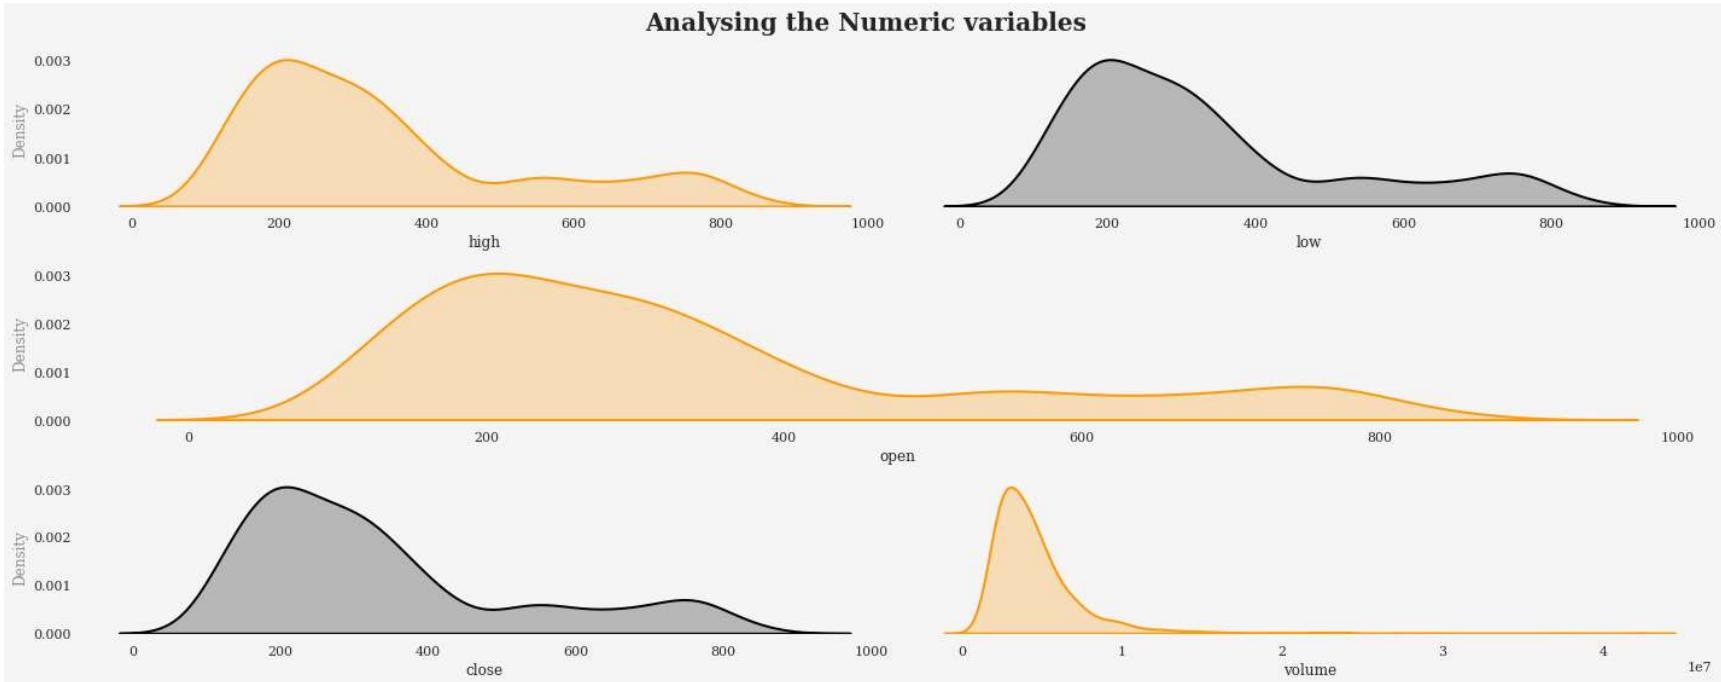
## Univariated Analysis

```
In [18]: #integer columns
fig=plt.figure(figsize=(20,8), tight_layout=True)
plt.suptitle("Analysing the Numeric variables", size=20, weight='bold')
ax=fig.subplot_mosaic("""AB
                      CC
                      DE""")
sns.kdeplot(df1['high'], ax=ax['A'], color=colors[0], fill=True, linewidth=2)
sns.kdeplot(df1['low'], ax=ax['B'], color=colors[1], fill=True, linewidth=2)
sns.kdeplot(df1['open'], ax=ax['C'], color=colors[0], fill=True, linewidth=2)
sns.kdeplot(df1['close'], ax=ax['D'], color=colors[1], fill=True, linewidth=2)
sns.kdeplot(df1['volume'], ax=ax['E'], color=colors[0], fill=True, linewidth=2)
ax['B'].yaxis.set_visible(False)
ax['E'].yaxis.set_visible(False)
ax['A'].yaxis.label.set_alpha(0.5)
ax['C'].yaxis.label.set_alpha(0.5)
ax['A'].yaxis.label.set_alpha(0.5)
ax['C'].yaxis.label.set_alpha(0.5)
```

```

ax['D'].yaxis.label.set_alpha(0.5)
for s in ['left','right','top','bottom']:
    ax['A'].spines[s].set_visible(False)
    ax['B'].spines[s].set_visible(False)
    ax['C'].spines[s].set_visible(False)
    ax['D'].spines[s].set_visible(False)
    ax['E'].spines[s].set_visible(False)

```



In [19]:

```

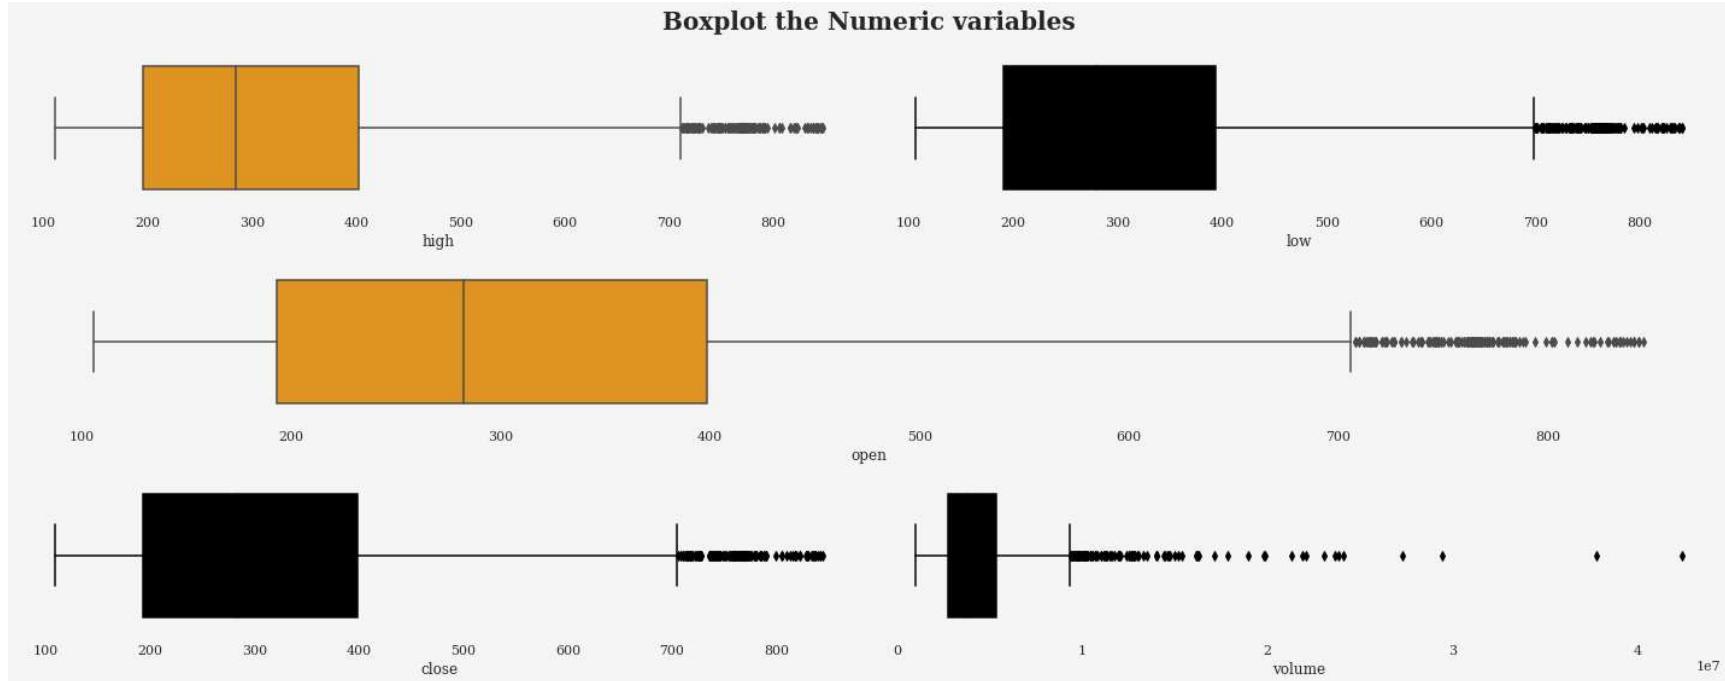
#integer columns
fig=plt.figure(figsize=(20,8), tight_layout=True)
plt.suptitle("Boxplot the Numeric variables", size=20, weight='bold')
ax=fig.subplot_mosaic('''
    AB
    CC
    DE''')
sns.boxplot(df1['high'], ax=ax['A'], color=colors[0])
sns.boxplot(df1['low'], ax=ax['B'], color=colors[1])
sns.boxplot(df1['open'], ax=ax['C'], color=colors[0])
sns.boxplot(df1['close'], ax=ax['D'], color=colors[1])
sns.boxplot(df1['volume'], ax=ax['E'], color=colors[1])
ax['B'].yaxis.set_visible(False)
ax['E'].yaxis.set_visible(False)
ax['A'].yaxis.label.set_alpha(0.5)
ax['C'].yaxis.label.set_alpha(0.5)

```

```

ax['A'].yaxis.label.set_alpha(0.5)
ax['C'].yaxis.label.set_alpha(0.5)
ax['D'].yaxis.label.set_alpha(0.5)
for s in ['left','right','top','bottom']:
    ax['A'].spines[s].set_visible(False)
    ax['B'].spines[s].set_visible(False)
    ax['C'].spines[s].set_visible(False)
    ax['D'].spines[s].set_visible(False)
    ax['E'].spines[s].set_visible(False)

```



*There seems to be outliers in the data set, however, those values can't be considered as outlier as they may be extrem values during peak selling days*

*Date column has been ignored as it is series of numbers*

## Categorical feature analysis

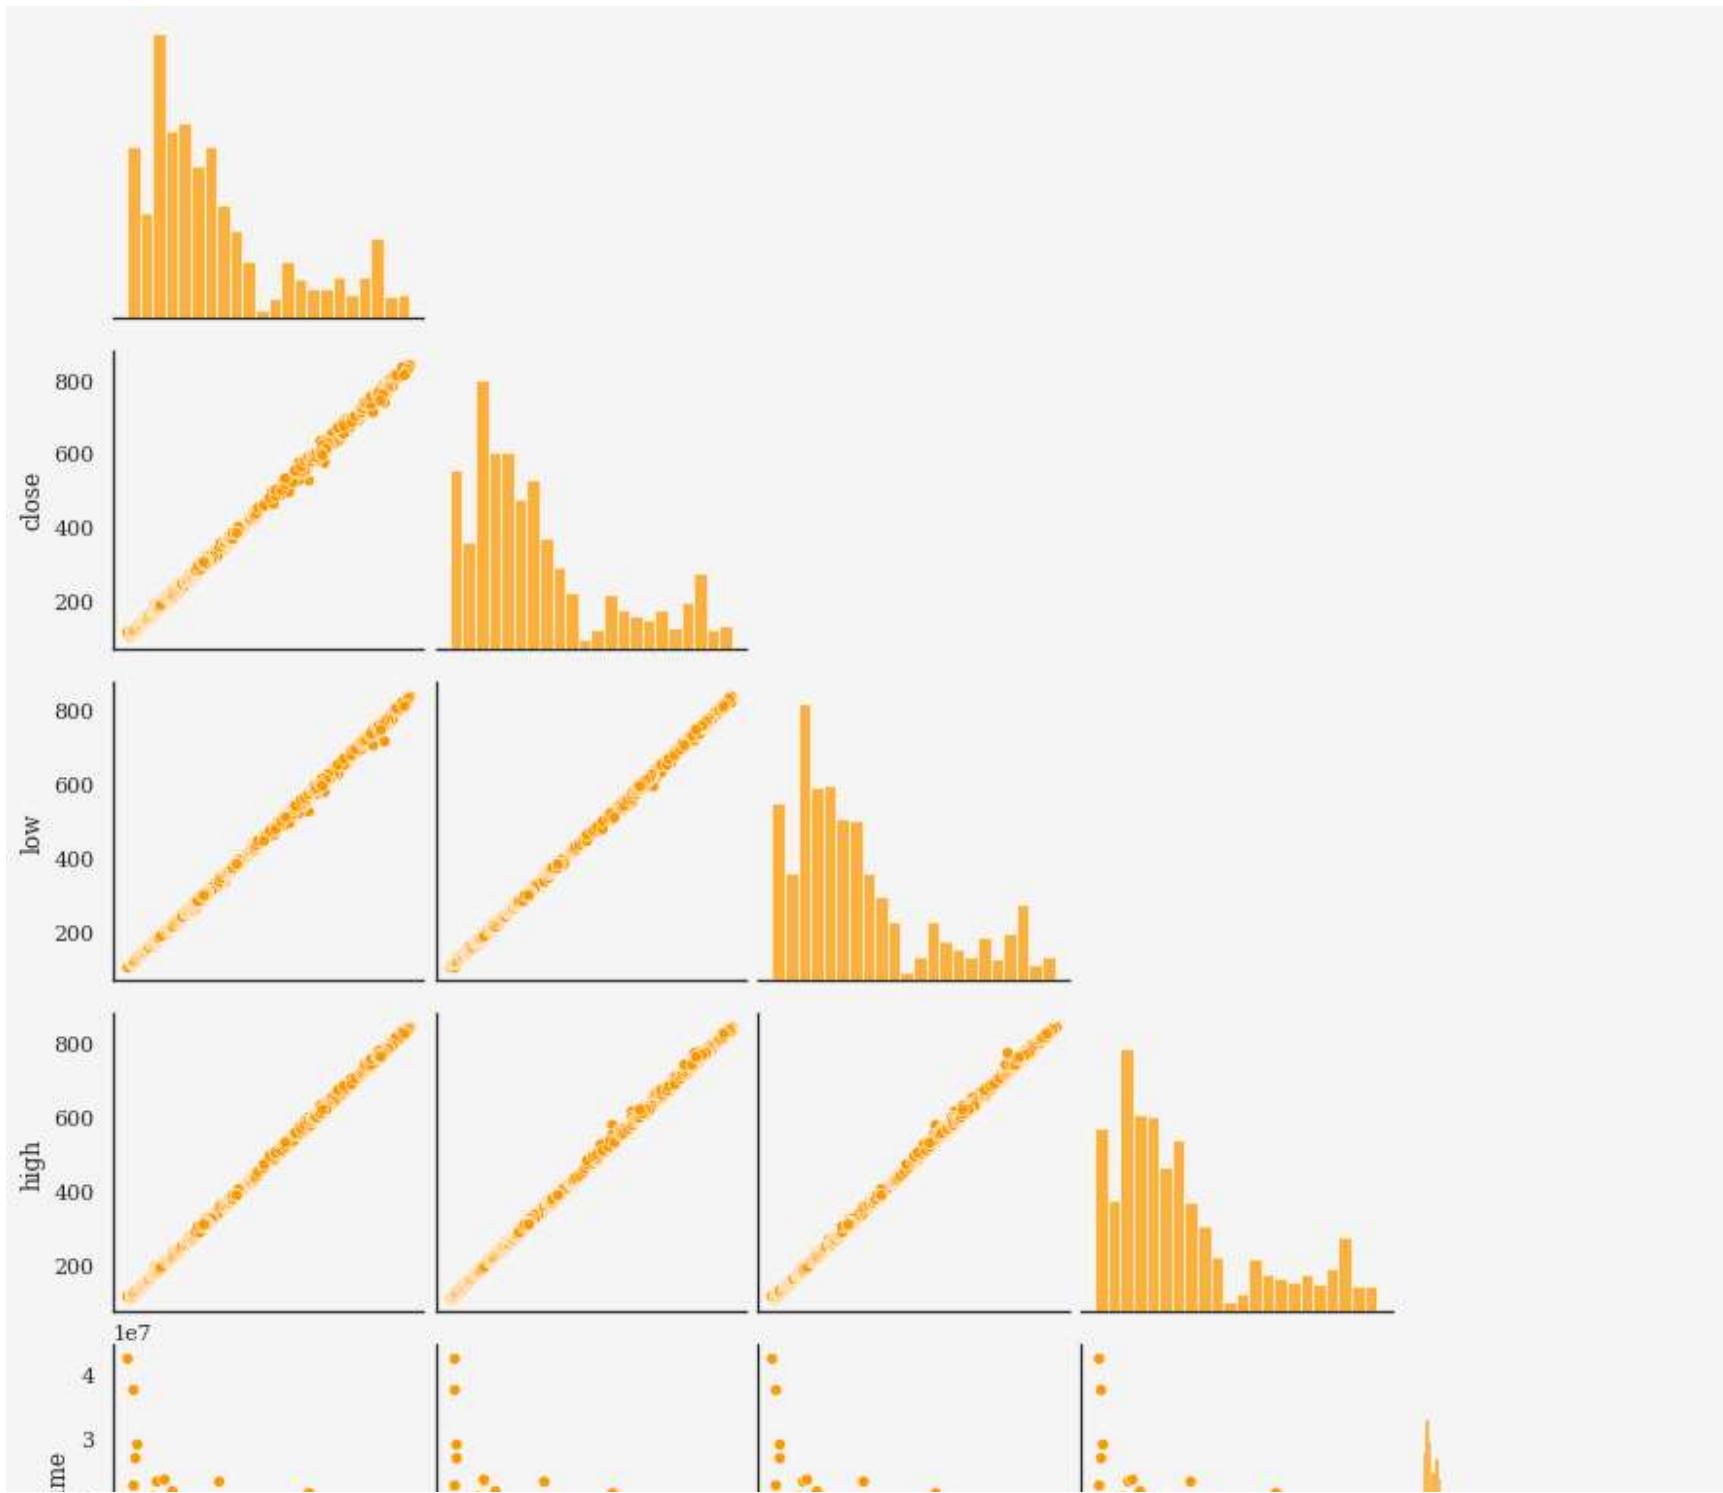
\*\* There is no categorical feature other than Symbol. as we have taken the Google stock as sample data, we will drop the Symbol feature.

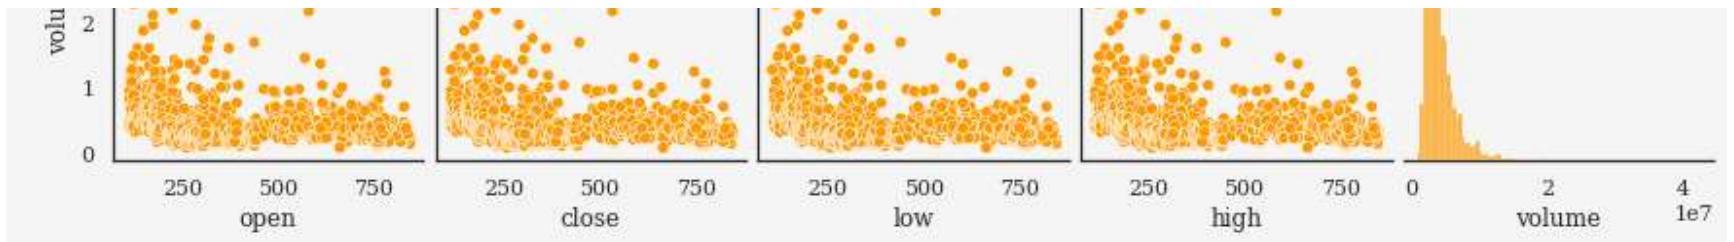
```
In [20]: df1.drop(['symbol'], axis=1, inplace=True)
```

## Bivariated & Multivariated Analysis

```
In [21]: # we need to predict the closing price of the stock, lets us consider 'Close' feature as the Target variable.  
sns.pairplot(df1,corner=True)
```

```
Out[21]: <seaborn.axisgrid.PairGrid at 0x79b7d80bc690>
```





```
In [22]: df1.corr()['close']
```

```
Out[22]: open      0.999581
          close     1.000000
          low       0.999832
          high      0.999811
          volume    -0.238560
Name: close, dtype: float64
```

## Hypothesis test to find the Normality in the Dataset

```
In [23]: from scipy.stats import levene, shapiro
int_cols=df1.select_dtypes(exclude='object').columns.to_list()

for i in int_cols:
    _, p_value=shapiro(df1[i])
    if p_value<0.05:
        print("Feature {} is normally distributed".format(i))
    else:
        print("Feature {} is not normaly distributed".format(i))

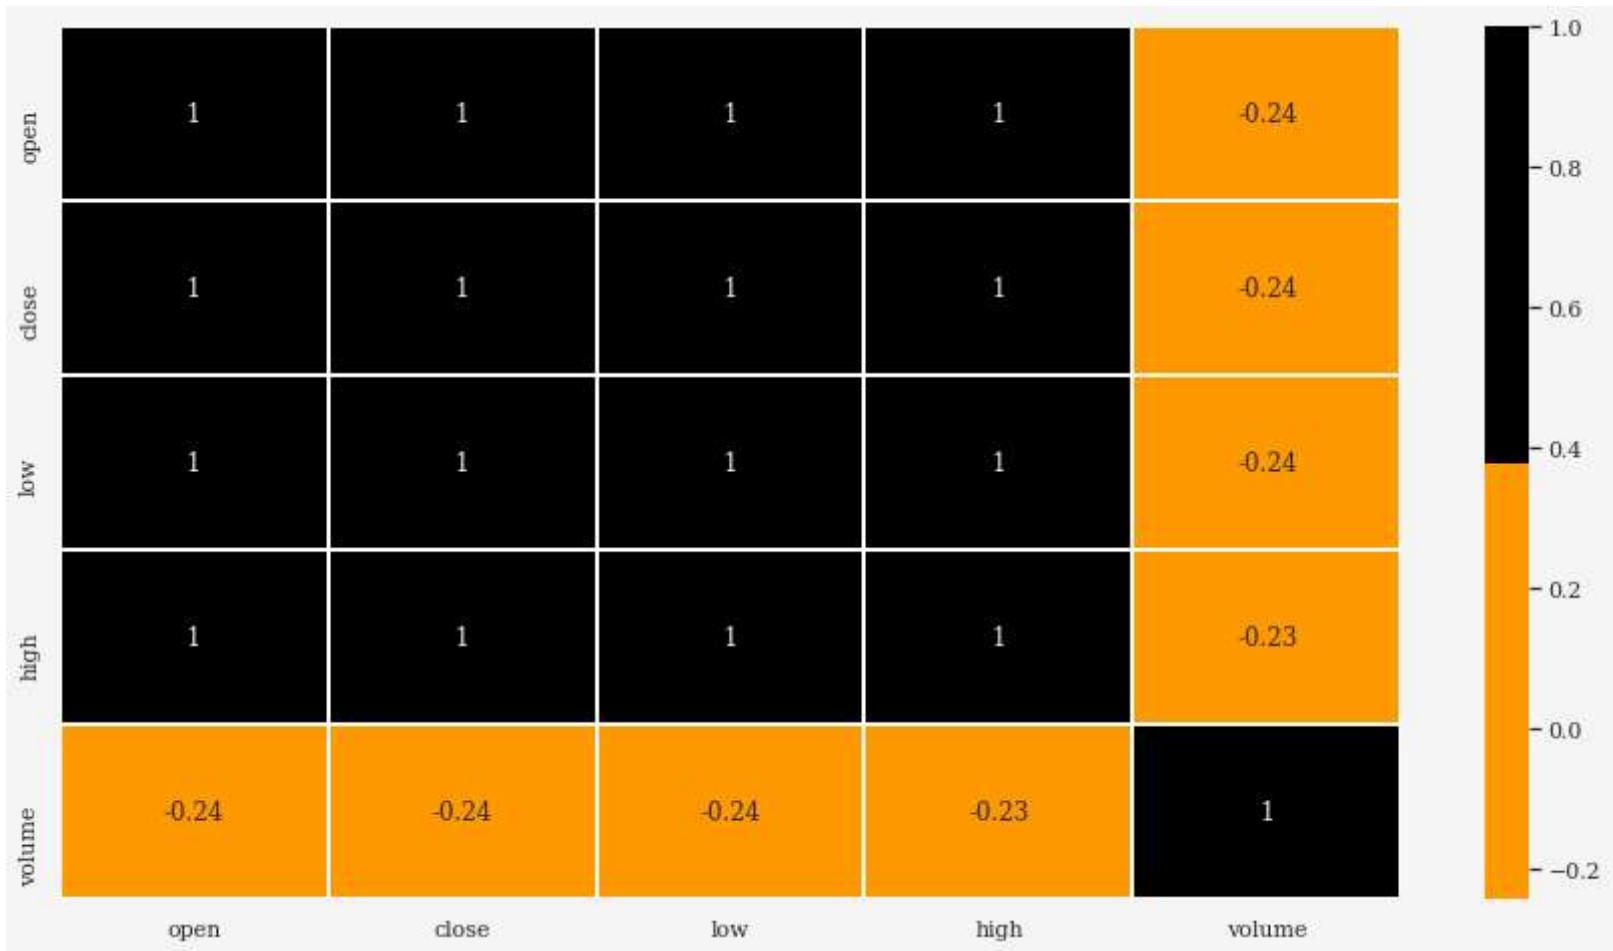
print("Normalitiy test p_value for feature -  {} is {}".format(i,np.round(p_value,3)))
```

```
Feature date is normaly distributed
Normalitiy test p_value for featue -  date is 0.0
Feature open is normaly distributed
Normalitiy test p_value for featue -  open is 0.0
Feature close is normaly distributed
Normalitiy test p_value for featue -  close is 0.0
Feature low is normaly distributed
Normalitiy test p_value for featue -  low is 0.0
Feature high is normaly distributed
Normalitiy test p_value for featue -  high is 0.0
Feature volume is normaly distributed
Normalitiy test p_value for featue -  volume is 0.0
```

## Correlation

```
In [24]: fig=plt.figure(figsize=(15,8))
sns.heatmap(df1.corr(), annot=True, cmap=[ colors[0],colors[1]], linecolor='white', linewidth=2 )
```

```
Out[24]: <AxesSubplot:>
```



**The feature Open, High, Low are highly correlated to Target feature Close. we can use either one of the feature for prediction to avoid multicollinearity**

## Train Test Split

```
In [25]: X=df1[['volume','open']]
y=df1['close']
```

```
In [26]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.3, shuffle=False, random_state=42)
```

# Normalizing the values

```
In [27]: from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```

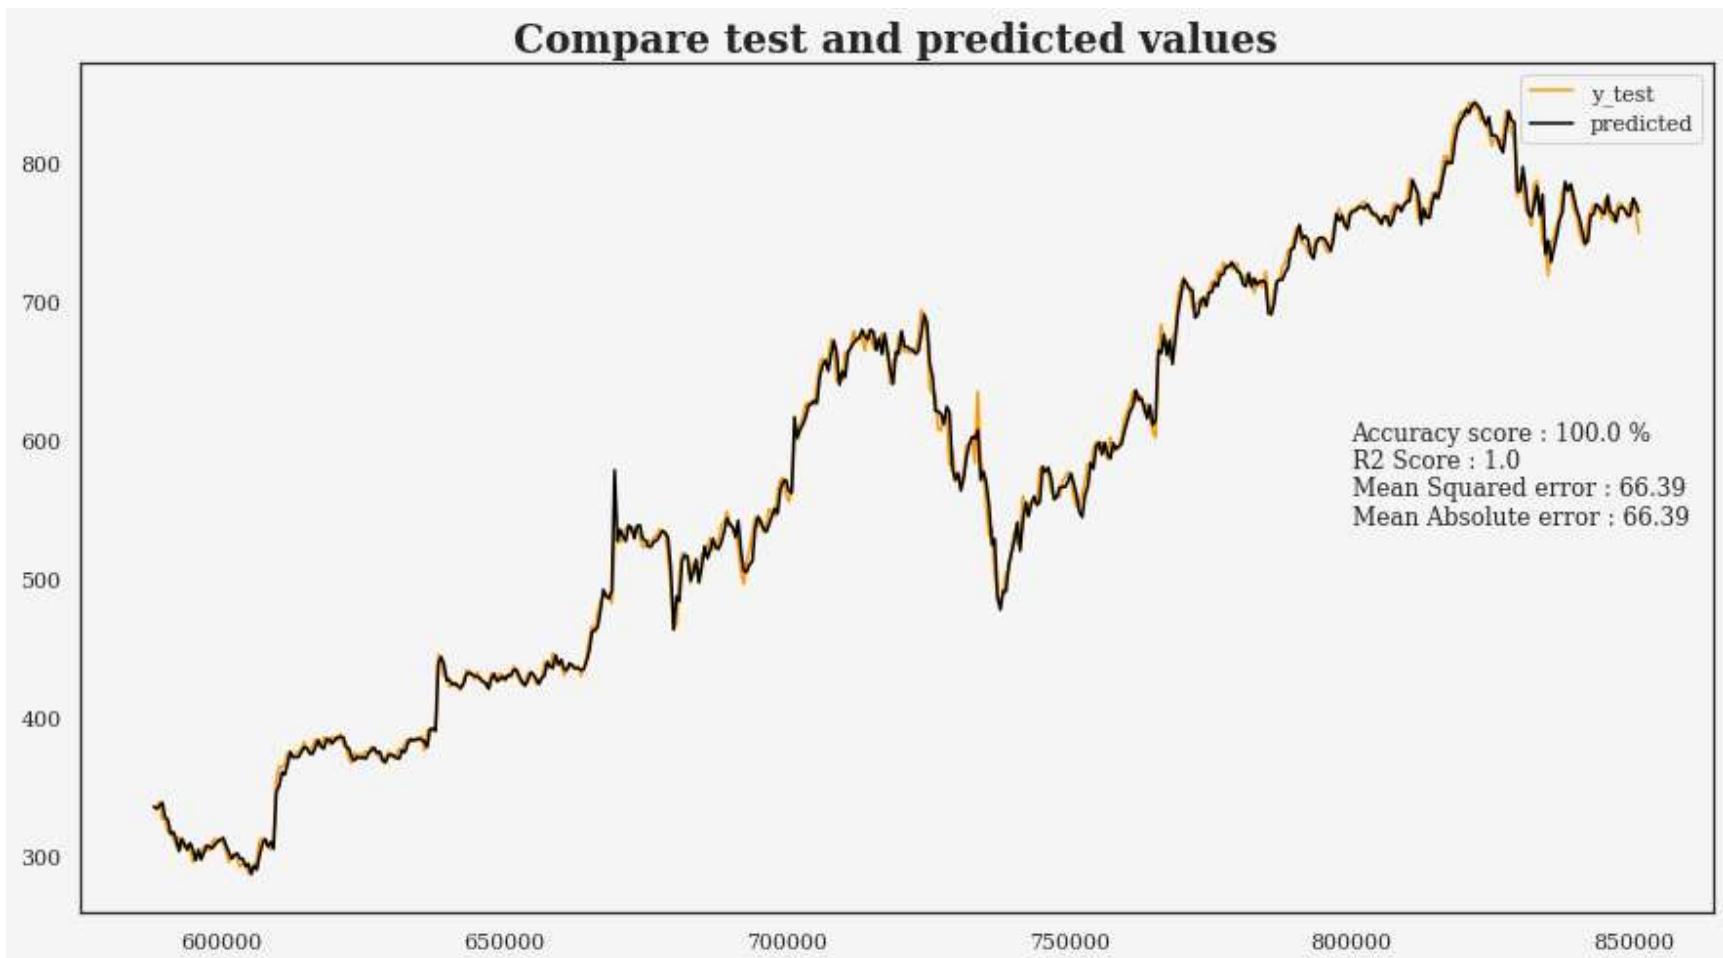
## Model Creation

### Linear Regression Model

```
In [28]: from sklearn.linear_model import LinearRegression  
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error  
from sklearn import set_config  
model = LinearRegression()  
model.fit(X_train,y_train)  
set_config(display='diagram')  
pred=model.predict(X_test)  
sc=np.round(model.score(X_test, y_test),2) * 100  
r2=np.round(r2_score(y_test,pred),2)  
mse=np.round(mean_squared_error(y_test,pred),2)  
mae=np.round(mean_squared_error(y_test,pred),2)
```

```
In [29]: fig=plt.figure(figsize=(15,8))  
p=pd.Series(pred, index=y_test.index)  
plt.plot(y_test)  
plt.plot(p)  
plt.legend(['y_test','predicted'])  
plt.title("Compare test and predicted values", size=20, weight='bold')  
plt.text(x=800000, y=600,s='Accuracy score : {}'.format(sc))  
plt.text(x=800000, y=580,s='R2 Score : {}'.format(r2))  
plt.text(x=800000, y=560,s='Mean Squared error : {}'.format(mse))  
plt.text(x=800000, y=540,s='Mean Absolute error : {}'.format(mae))
```

```
Out[29]: Text(800000, 540, 'Mean Absolute error : 66.39')
```



## LSTM Model

```
In [30]: X=df1[['open','high']]
y=df1['close']
length=100
#from sklearn.model_selection import train_test_split
#X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2, shuffle=False, random_state=42)
training_set = X.iloc[:1000].values
test_set = X.iloc[1000:1000+length].values
# Feature Scaling
from sklearn.preprocessing import MinMaxScaler
sc = MinMaxScaler(feature_range = (0, 1))
```

```

training_set_scaled = sc.fit_transform(training_set)
test_set_scaled=sc.transform(test_set)
# Creating a data structure with 60 time-steps and 1 output
X_train = []
y_train = []
for i in range(length, len(training_set)):
    X_train.append(training_set_scaled[i-length:i, 0])
    y_train.append(training_set_scaled[i, 0])
X_train, y_train = np.array(X_train), np.array(y_train)
X_train = np.reshape(X_train, (X_train.shape[0], X_train.shape[1], 1))

X_test = []
y_test = []
for i in range(length, len(test_set)):
    X_test.append(test_set_scaled[i-length:i, 0])
    y_test.append(test_set_scaled[i, 0])
X_test, y_test = np.array(X_test), np.array(y_test)
X_test = np.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))

```

In [31]:

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM, Dropout
model = Sequential()
#Adding the first LSTM Layer and some Dropout regularisation
model.add(LSTM(units = 50, return_sequences = True, input_shape = (X_train.shape[1], 1)))
model.add(Dropout(0.2))
# Adding a second LSTM Layer and some Dropout regularisation
model.add(LSTM(units = 50, return_sequences = True))
model.add(Dropout(0.2))
# Adding a third LSTM Layer and some Dropout regularisation
model.add(LSTM(units = 50, return_sequences = True))
model.add(Dropout(0.2))
# Adding a fourth LSTM Layer and some Dropout regularisation
model.add(LSTM(units = 50))
model.add(Dropout(0.2))
# Adding the output layer
model.add(Dense(units = 1))

# Compiling the RNN
model.compile(optimizer = 'adam', loss = 'mean_squared_error')

# Fitting the RNN to the Training set
model.fit(X_train, y_train, validation_data=(X_test,y_test),epochs = 100, batch_size = 32)

```

Epoch 1/100  
29/29 [=====] - 9s 74ms/step - loss: 0.0613 - val\_loss: 0.4825  
Epoch 2/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0080 - val\_loss: 0.1162  
Epoch 3/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0040 - val\_loss: 0.1090  
Epoch 4/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0040 - val\_loss: 0.0377  
Epoch 5/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0039 - val\_loss: 0.0403  
Epoch 6/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0034 - val\_loss: 0.0694  
Epoch 7/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0037 - val\_loss: 0.0645  
Epoch 8/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0039 - val\_loss: 0.0570  
Epoch 9/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0029 - val\_loss: 0.0596  
Epoch 10/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0036 - val\_loss: 0.0580  
Epoch 11/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0029 - val\_loss: 0.0228  
Epoch 12/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0032 - val\_loss: 0.0442  
Epoch 13/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0035 - val\_loss: 0.0441  
Epoch 14/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0027 - val\_loss: 0.0700  
Epoch 15/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0030 - val\_loss: 0.0410  
Epoch 16/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0027 - val\_loss: 0.0328  
Epoch 17/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0029 - val\_loss: 0.0363  
Epoch 18/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0025 - val\_loss: 0.0281  
Epoch 19/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0027 - val\_loss: 0.0161  
Epoch 20/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0027 - val\_loss: 0.0171  
Epoch 21/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0024 - val\_loss: 0.0224  
Epoch 22/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0023 - val\_loss: 0.0156  
Epoch 23/100

```
29/29 [=====] - 1s 24ms/step - loss: 0.0027 - val_loss: 0.0396
Epoch 24/100
29/29 [=====] - 1s 24ms/step - loss: 0.0026 - val_loss: 0.0277
Epoch 25/100
29/29 [=====] - 1s 24ms/step - loss: 0.0025 - val_loss: 0.0238
Epoch 26/100
29/29 [=====] - 1s 24ms/step - loss: 0.0024 - val_loss: 0.0129
Epoch 27/100
29/29 [=====] - 1s 24ms/step - loss: 0.0025 - val_loss: 0.0124
Epoch 28/100
29/29 [=====] - 1s 24ms/step - loss: 0.0023 - val_loss: 0.0148
Epoch 29/100
29/29 [=====] - 1s 24ms/step - loss: 0.0025 - val_loss: 0.0272
Epoch 30/100
29/29 [=====] - 1s 24ms/step - loss: 0.0020 - val_loss: 0.0224
Epoch 31/100
29/29 [=====] - 1s 25ms/step - loss: 0.0018 - val_loss: 0.0312
Epoch 32/100
29/29 [=====] - 1s 24ms/step - loss: 0.0021 - val_loss: 0.0132
Epoch 33/100
29/29 [=====] - 1s 24ms/step - loss: 0.0028 - val_loss: 0.0088
Epoch 34/100
29/29 [=====] - 1s 24ms/step - loss: 0.0022 - val_loss: 0.0130
Epoch 35/100
29/29 [=====] - 1s 24ms/step - loss: 0.0023 - val_loss: 0.0096
Epoch 36/100
29/29 [=====] - 1s 24ms/step - loss: 0.0022 - val_loss: 0.0094
Epoch 37/100
29/29 [=====] - 1s 24ms/step - loss: 0.0021 - val_loss: 0.0174
Epoch 38/100
29/29 [=====] - 1s 24ms/step - loss: 0.0019 - val_loss: 0.0165
Epoch 39/100
29/29 [=====] - 1s 24ms/step - loss: 0.0022 - val_loss: 0.0168
Epoch 40/100
29/29 [=====] - 1s 24ms/step - loss: 0.0020 - val_loss: 0.0106
Epoch 41/100
29/29 [=====] - 1s 24ms/step - loss: 0.0022 - val_loss: 0.0108
Epoch 42/100
29/29 [=====] - 1s 24ms/step - loss: 0.0020 - val_loss: 0.0142
Epoch 43/100
29/29 [=====] - 1s 24ms/step - loss: 0.0019 - val_loss: 0.0127
Epoch 44/100
29/29 [=====] - 1s 24ms/step - loss: 0.0019 - val_loss: 0.0104
Epoch 45/100
29/29 [=====] - 1s 24ms/step - loss: 0.0020 - val_loss: 0.0077
```

Epoch 46/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0018 - val\_loss: 0.0184  
Epoch 47/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0019 - val\_loss: 0.0126  
Epoch 48/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0016 - val\_loss: 0.0089  
Epoch 49/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0016 - val\_loss: 0.0143  
Epoch 50/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0019 - val\_loss: 0.0098  
Epoch 51/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0017 - val\_loss: 0.0266  
Epoch 52/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0018 - val\_loss: 0.0138  
Epoch 53/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0018 - val\_loss: 0.0102  
Epoch 54/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0017 - val\_loss: 0.0077  
Epoch 55/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0028 - val\_loss: 0.0140  
Epoch 56/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0015 - val\_loss: 0.0180  
Epoch 57/100  
29/29 [=====] - 1s 25ms/step - loss: 0.0016 - val\_loss: 0.0124  
Epoch 58/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0012 - val\_loss: 0.0067  
Epoch 59/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0015 - val\_loss: 0.0100  
Epoch 60/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0013 - val\_loss: 0.0135  
Epoch 61/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0015 - val\_loss: 0.0087  
Epoch 62/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0014 - val\_loss: 0.0065  
Epoch 63/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0015 - val\_loss: 0.0175  
Epoch 64/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0015 - val\_loss: 0.0137  
Epoch 65/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0014 - val\_loss: 0.0098  
Epoch 66/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0012 - val\_loss: 0.0119  
Epoch 67/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0015 - val\_loss: 0.0072  
Epoch 68/100

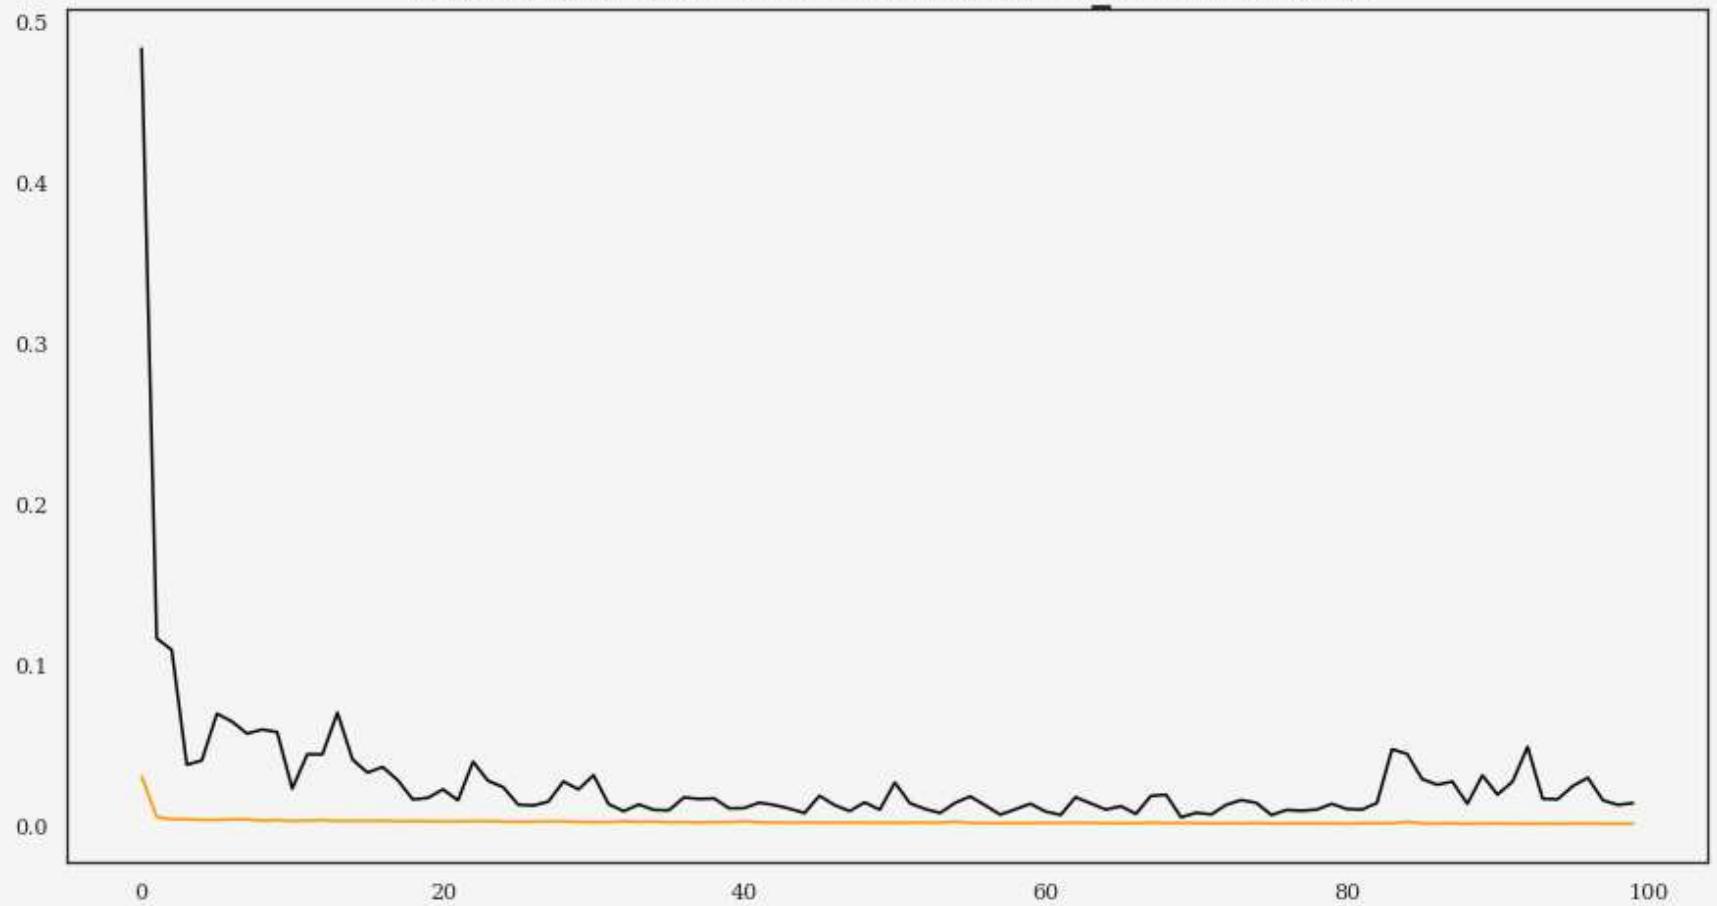
29/29 [=====] - 1s 24ms/step - loss: 0.0018 - val\_loss: 0.0184  
Epoch 69/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0015 - val\_loss: 0.0191  
Epoch 70/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0014 - val\_loss: 0.0051  
Epoch 71/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0015 - val\_loss: 0.0078  
Epoch 72/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0012 - val\_loss: 0.0068  
Epoch 73/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0016 - val\_loss: 0.0128  
Epoch 74/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0012 - val\_loss: 0.0158  
Epoch 75/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0015 - val\_loss: 0.0140  
Epoch 76/100  
29/29 [=====] - 1s 25ms/step - loss: 0.0013 - val\_loss: 0.0064  
Epoch 77/100  
29/29 [=====] - 1s 25ms/step - loss: 0.0012 - val\_loss: 0.0096  
Epoch 78/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0013 - val\_loss: 0.0091  
Epoch 79/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0012 - val\_loss: 0.0098  
Epoch 80/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0012 - val\_loss: 0.0134  
Epoch 81/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0011 - val\_loss: 0.0102  
Epoch 82/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0012 - val\_loss: 0.0098  
Epoch 83/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0013 - val\_loss: 0.0140  
Epoch 84/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0012 - val\_loss: 0.0473  
Epoch 85/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0027 - val\_loss: 0.0444  
Epoch 86/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0014 - val\_loss: 0.0288  
Epoch 87/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0011 - val\_loss: 0.0253  
Epoch 88/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0013 - val\_loss: 0.0272  
Epoch 89/100  
29/29 [=====] - 1s 24ms/step - loss: 9.8388e-04 - val\_loss: 0.0134  
Epoch 90/100  
29/29 [=====] - 1s 24ms/step - loss: 0.0011 - val\_loss: 0.0311

```
Epoch 91/100
29/29 [=====] - 1s 24ms/step - loss: 0.0014 - val_loss: 0.0192
Epoch 92/100
29/29 [=====] - 1s 24ms/step - loss: 0.0010 - val_loss: 0.0272
Epoch 93/100
29/29 [=====] - 1s 24ms/step - loss: 0.0013 - val_loss: 0.0489
Epoch 94/100
29/29 [=====] - 1s 24ms/step - loss: 0.0012 - val_loss: 0.0166
Epoch 95/100
29/29 [=====] - 1s 24ms/step - loss: 9.8453e-04 - val_loss: 0.0162
Epoch 96/100
29/29 [=====] - 1s 24ms/step - loss: 0.0012 - val_loss: 0.0245
Epoch 97/100
29/29 [=====] - 1s 24ms/step - loss: 0.0011 - val_loss: 0.0296
Epoch 98/100
29/29 [=====] - 1s 24ms/step - loss: 9.5980e-04 - val_loss: 0.0156
Epoch 99/100
29/29 [=====] - 1s 24ms/step - loss: 9.3465e-04 - val_loss: 0.0128
Epoch 100/100
29/29 [=====] - 1s 24ms/step - loss: 0.0011 - val_loss: 0.0139
Out[31]: <tensorflow.python.keras.callbacks.History at 0x79b7981594d0>
```

```
In [32]: loss = pd.DataFrame(model.history.history)
fig=plt.figure(figsize=(15,8))
plt.title("Validation loss Vs Train inverse_transforms", size=20, weight='bold')
plt.plot(loss)
```

```
Out[32]: [<matplotlib.lines.Line2D at 0x79b72819b590>,
<matplotlib.lines.Line2D at 0x79b728134ad0>]
```

## Validation loss Vs Train inverse\_transforms



```
In [33]: pred=model.predict(X_test)
test=pd.DataFrame(columns=['test','pred'])
test['test']=y_test
test['pred']=pred.flatten()
test
```

Out[33]:

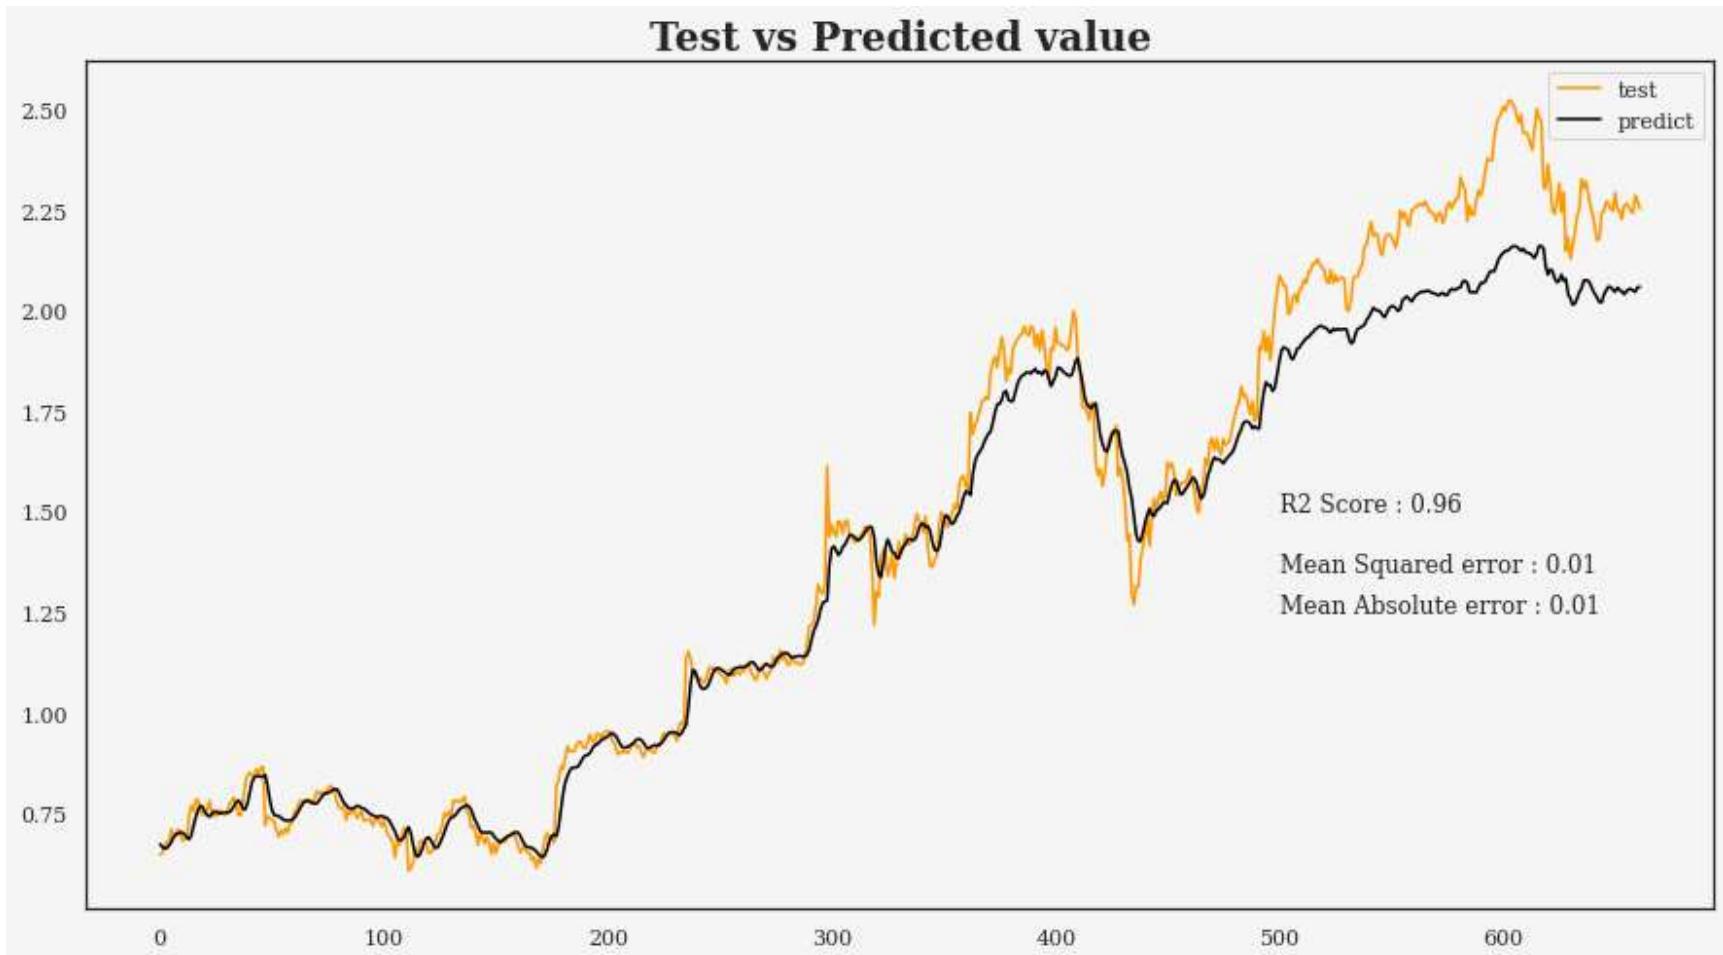
	test	pred
0	0.647729	0.674112
1	0.652302	0.667890
2	0.669738	0.663177
3	0.679428	0.664529
4	0.680827	0.671369
...	...	...
657	2.247313	2.054365
658	2.243389	2.050452
659	2.287235	2.047239
660	2.274098	2.058147
661	2.253864	2.059852

662 rows × 2 columns

In [34]:

```
fig=plt.figure(figsize=(15,8))
plt.title("Test vs Predicted value", size=20, weight='bold')
plt.plot(test)
plt.legend(['test','predict'])
r2=np.round(r2_score(y_test,pred),2)
mse=np.round(mean_squared_error(y_test,pred),2)
mae=np.round(mean_squared_error(y_test,pred),2)
plt.text(x=500, y=1.5,s='R2 Score : {}'.format(r2))
plt.text(x=500, y=1.35,s='Mean Squared error : {}'.format(mse))
plt.text(x=500, y=1.25,s='Mean Absolute error : {}'.format(mae))
```

Out[34]: Text(500, 1.25, 'Mean Absolute error : 0.01')



## FaceBook Prophet Model

```
In [35]: from fbprophet import Prophet
df_p = df1[['date','close']]
df_p.columns=['ds','y']

split_data = df_p.index.max()-100000
train = df_p.loc[df_p.index<=split_data].copy()
test=df_p.loc[df_p.index>split_data].copy()
train.set_index('ds',inplace=True)
test.set_index('ds',inplace=True)
train.reset_index(inplace=True)
```

```
test.reset_index(inplace=True)

#Model creation
model=Prophet()
model.fit(train)

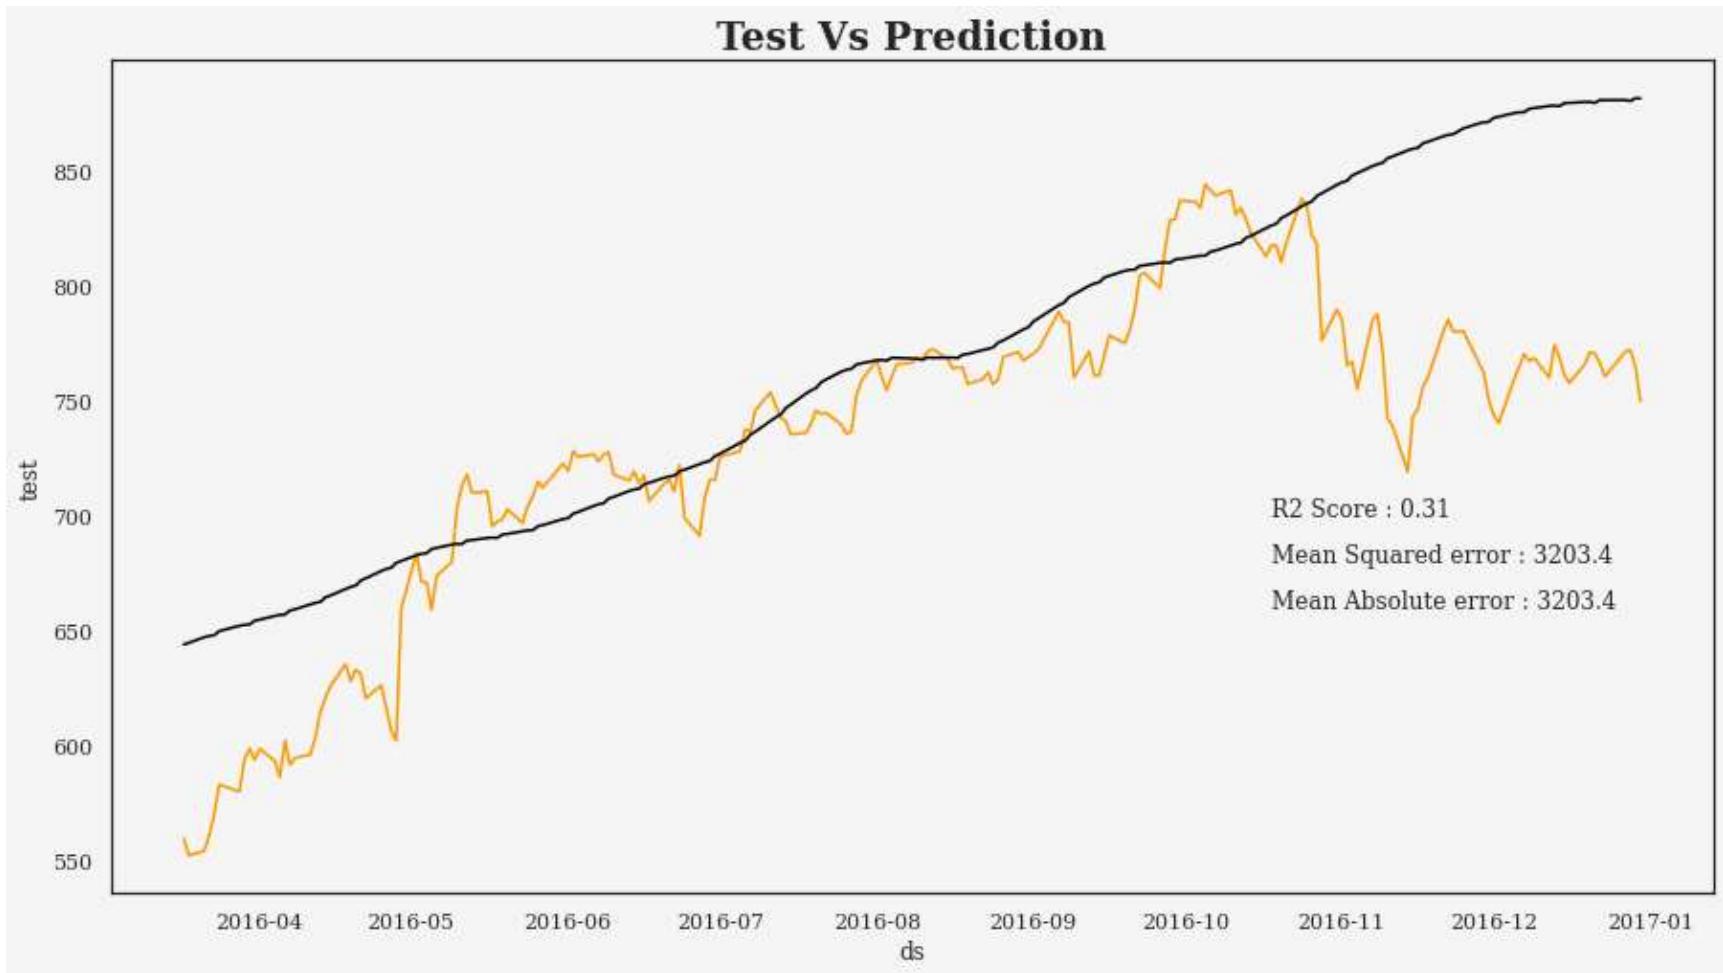
#model prediction
pred=model.predict(test)
```

```
In [36]: test_pred=pd.DataFrame(columns=['ds','test','predict','pred_lower','pred_high'], index=test.index)
test_pred['test']=test['y']
test_pred['ds']=test['ds']
test_pred['predict']=pred['yhat']
test_pred['pred_lower']=pred['yhat_lower']
test_pred['pred_high']=pred['yhat_upper']
```

```
In [37]: #plotting Test vs Predicted
fig=plt.figure(figsize=(15,8))
plt.title("Test Vs Prediction", size=20, weight='bold')
sns.lineplot(data=test_pred,x='ds',y='test')
sns.lineplot(data=test_pred,x='ds',y='predict')

r2=np.round(r2_score(test_pred['test'],test_pred['predict']),2)
mse=np.round(mean_squared_error(test_pred['test'],test_pred['predict']),2)
mae=np.round(mean_squared_error(test_pred['test'],test_pred['predict']),2)
plt.text(x=mdates.datestr2num('2016-10'), y=700,s='R2 Score : {}'.format(r2))
plt.text(x=mdates.datestr2num('2016-10'), y=680,s='Mean Squared error : {}'.format(mse))
plt.text(x=mdates.datestr2num('2016-10'), y=660,s='Mean Absolute error : {}'.format(mae))
```

```
Out[37]: Text(17092.0, 660, 'Mean Absolute error : 3203.4')
```



```
In [38]: fig=plt.figure(figsize=(15,8))
plt.title("Test Vs Prediction", size=20, weight='bold')
sns.lineplot(data=train,x='ds',y='y')
sns.lineplot(data=test_pred,x='ds',y='predict')
sns.lineplot(data=test_pred,x='ds',y='test', alpha=0.5, ls='--', color='black')

r2=np.round(r2_score(test_pred['test'],test_pred['predict']),2)
mse=np.round(mean_squared_error(test_pred['test'],test_pred['predict']),2)
mae=np.round(mean_squared_error(test_pred['test'],test_pred['predict']),2)
plt.text(x=mdates.datestr2num('2015'), y=300,s='R2 Score : {}'.format(r2))
plt.text(x=mdates.datestr2num('2015'), y=260,s='Mean Squared error : {}'.format(mse))
plt.text(x=mdates.datestr2num('2015'), y=220,s='Mean Absolute error : {}'.format(mae))
plt.legend(['Train','Predict','Test'])
```

Out[38]: <matplotlib.legend.Legend at 0x79b30607c590>

